

✓ Advanced Recommender Systems 2

Agenda

1. Content-Based Filtering

- Key Principles of CBF
- Steps in Content Based Filtering
- Feature Representation
- Similarity Metrics
- Term Frequency-Inverse Document Frequency

2. Hybrid Recommender Systems

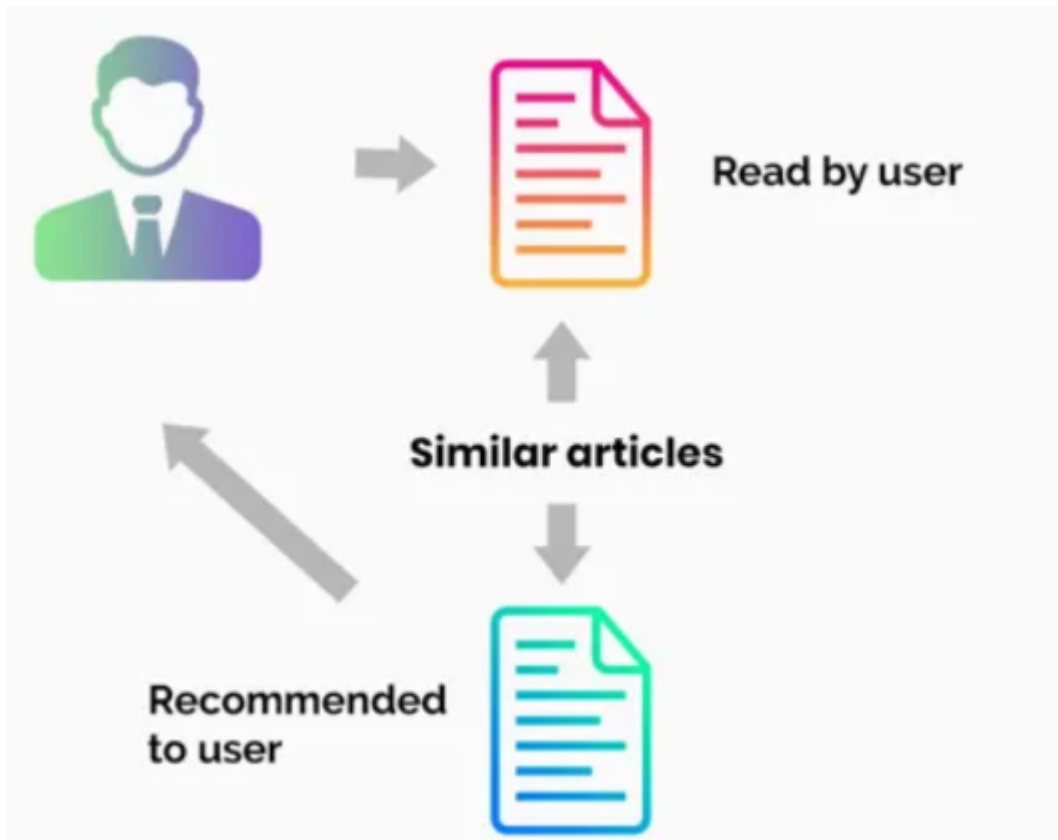
- Key Components of HRS
- Types of Hybrid recommender Systems

3. Handling Cold Start Problems

- Types of Cold Start Problems
- Challenges in Cold Start Scenarios
- Strategies for Handling Cold Start Problems

✓ Content-Based Filtering

- Content-Based Filtering (CBF) is a recommender system technique that suggests items by analyzing the attributes of the items themselves and the user's preferences.
- It builds a user profile by tracking their past interactions, likes, and preferences, focusing on the features or characteristics of the items they have shown interest in.
- The system then recommends new items with similar attributes to those previously favored by the user.
- For example, if a user enjoys action movies, CBF would recommend other action films based on genre, actors, or directors.
- Unlike collaborative filtering, CBF doesn't rely on other users' behavior, making it effective even when data on other users is limited.



✓ Key Principles of Content-Based Filtering

- **Item Profiles:** Each item in the system is described using a set of features or attributes. For example, a movie can be described by its genre, director, cast, language, and other metadata. These attributes are used to create a feature vector that represents the item.
- **User Profiles:** Each user is represented by a profile, which is a collection of the user's preferences, tastes, or history of interactions. The user profile can be built by analyzing the items the user has rated, liked, or consumed in the past. The system matches these items with the user's preferences to suggest similar items.
- **Similarity Calculation:** The recommender system calculates the similarity between the feature vectors of items the user has liked and other items in the database. This similarity is usually measured using a metric such as cosine similarity, Euclidean distance, or Pearson correlation.
- **Personalized Recommendations:** Based on the similarity score, the system ranks the items and recommends those that are most similar to the items the user has already interacted with.

✓ Steps in Content-Based Filtering

1. Feature Extraction:

- The first step in CBF is to extract relevant features from the items.
- In domains like movies, books, or articles, these features can include textual attributes (e.g., keywords, genres), categorical attributes (e.g., type, category), and sometimes numeric attributes (e.g., duration or length).
- These extracted features form the basis for comparison with user preferences.

2. User Profile Creation:

- Next, a user profile is created based on their past interactions, such as items they've rated, liked, or consumed.
- The profile may be weighted to give more importance to recent or higher-rated items.
- For instance, if a user consistently rates action movies highly, their profile will reflect a strong preference for the "action" genre.

3. Utility matrix:

- A utility matrix represents the interactions between users and items, capturing how users engage with different items based on their preferences.
- This matrix stores data collected from users' daily activities, such as ratings, purchases, or views, and organizes it in a structured format to help identify the user's likes and dislikes.
- Each interaction in the matrix is assigned a numerical value, referred to as the 'degree of preference,' indicating how much a user likes or prefers a particular item.

4. Similarity Computation:

- To match user preferences with available items, the system computes the similarity between the user profile and the item feature vectors.
- Techniques like cosine similarity or dot product are commonly used to measure how closely an item's features align with the user's interests.

5. Recommendation Generation:

- Finally, based on the computed similarity scores, the system recommends the top-ranked items that most closely align with the user's preferences, ensuring personalized and relevant recommendations.

✓ Feature Representation

The success of content-based filtering (CBF) hinges on the quality and relevance of the features used to represent items. The feature vectors for items can consist of various types of attributes:

- **Text-based Features:** These include terms extracted from item descriptions or metadata, often represented using techniques like TF-IDF (Term Frequency-Inverse Document Frequency). For example, documents can be represented by the key words or phrases in their text.

- **Categorical Features:** Tags, genres, or categories associated with items. For instance, a book could belong to categories like "science fiction" or "romance," and movies might be categorized by genres like "comedy" or "drama."
- **Numerical Features:** Numeric attributes such as price, rating, or duration. For example, a product's price or a movie's duration can be included as features to guide recommendations.
- **Multimedia Features:** In recommendations involving images, music, or video, features could include pixel-level data for images or sound patterns in audio files. These features help capture the essence of multimedia content for personalized suggestions.

✓ Similarity Metrics

1. Cosine Similarity:

- Cosine similarity measures the cosine of the angle between two vectors in a multidimensional space, giving an indication of how similar they are.
- Formula:

$$\text{Cosine Similarity} = \frac{\sum_{i=1}^n A_i \cdot B_i}{\sqrt{\sum_{i=1}^n A_i^2} \cdot \sqrt{\sum_{i=1}^n B_i^2}}$$

- Values range between -1 (completely dissimilar) and 1 (perfectly similar).
- 1 indicates that the vectors point in the same direction, while -1 indicates they point in opposite directions.

This metric is commonly used in text and multimedia data, where the magnitude of the vectors doesn't matter as much as their direction.

2. Euclidean Distance:

- Euclidean distance is the straight-line distance between two vectors in a multidimensional space, where smaller distances indicate higher similarity.
- Formula:

$$\text{Euclidean Distance} = \sqrt{\sum_{i=1}^n (A_i - B_i)^2}$$

- The closer the distance is to zero, the more similar the vectors are.

- It's commonly used in cases where the magnitude of the features matters (e.g., continuous data).

3. Pearson Correlation:

Pearson correlation measures the linear correlation between two sets of data, and it's effective when the range of values varies significantly between items or users.

- It takes into account both the direction and strength of the linear relationship between two variables.
- Pearson correlation returns values between -1 (perfect negative correlation) and 1 (perfect positive correlation), with 0 indicating no correlation.

✓ Term Frequency-Inverse Document Frequency

- Term Frequency-Inverse Document Frequency (TF-IDF) is a widely used technique in natural language processing and text mining, and it can also be applied to content-based recommender systems.
- In the context of recommendation systems, particularly content-based filtering, TF-IDF helps to determine the relevance of terms (words or features) in a document or item and the importance of those terms across a larger collection of items.

What is TF-IDF?

TF-IDF combines two measures:

- Term Frequency (TF): It measures how often a term appears in a document (or item) compared to the total number of terms in the document.

$$TF(t, d) = \frac{\text{Number of times term } t \text{ appears in document } d}{\text{Total number of terms in document } d}$$

- Inverse Document Frequency (IDF): It measures the importance of a term across the entire corpus of documents (or items). If a term appears in many documents, it is considered less important.

$$IDF(t) = \log \left(\frac{\text{Total number of documents}}{\text{Number of documents containing term } t} \right)$$

- The final TF-IDF score is the product of the two:

$$TF - IDF(t, d) = TF(t, d) \times IDF(t)$$

This score reflects the importance of a term in a specific document while also considering how common or rare the term is across the entire collection.

✓ Applying TF-IDF in Recommender Systems

In content-based recommender systems, TF-IDF is typically applied to extract important features (words, tags, attributes) from the content associated with the items (e.g., descriptions, reviews, tags). Here's how it can be used:

- **Document Representation:**

- Items (e.g., books, movies, products) are treated as "documents" that contain content, such as text descriptions, reviews, or tags.
- Each item is represented as a vector of terms (keywords or features) using the TF-IDF scores. The higher the TF-IDF score of a term in an item, the more relevant it is to that item.

- **User Profile Representation:**

- A user profile is built by analyzing the content of items they have interacted with (e.g., rated, purchased, or viewed).
- TF-IDF can be used to represent the user's preferences based on the terms (features) of the items they liked. For instance, if a user has interacted with several "science fiction" books, the user profile will have high TF-IDF scores for terms like "sci-fi", "space", "future", etc.

- **Similarity Computation:**

- Cosine Similarity: To make recommendations, the system computes the cosine similarity between the TF-IDF vector of a user's profile and the TF-IDF vectors of other items in the system.
- Items with the highest similarity to the user's profile are recommended, as they share relevant terms (content) with the user's past interactions.

✓ Example: Content-Based Filtering Recommender System

- We will use the following steps in our implementation:
 - Create a toy dataset of movies.
 - Extract features from the dataset.
 - Calculate the cosine similarity between movies.
 - Recommend similar movies based on a user's preference.

```
pip install pandas scikit-learn
```

```
⇒ Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packag  
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-  
Requirement already satisfied: numpy>=1.22.4 in /usr/local/lib/python3.10/dis-  
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/pytho  
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-  
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/di-  
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.10/dist-  
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dis-  
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-pac
```

Creating a Toy Dataset

```

import pandas as pd

# Sample dataset of movies
data = {
    'title': [
        'The Matrix',
        'John Wick',
        'The Godfather',
        'Pulp Fiction',
        'The Dark Knight',
        'Inception',
        'The Lord of the Rings',
        'The Avengers'
    ],
    'genre': [
        'Action, Sci-Fi',
        'Action, Thriller',
        'Crime, Drama',
        'Crime, Drama',
        'Action, Crime',
        'Action, Sci-Fi',
        'Adventure, Fantasy',
        'Action, Sci-Fi'
    ],
    'description': [
        'A computer hacker learns about the true nature of reality and his role i
        'An ex-hitman comes out of retirement to track down the gangsters that ki
        'An organized crime dynasty\'s aging patriarch transfers control of his c
        'The lives of two mob hitmen, a boxer, a gangster`s wife, and a pair of d
        'When the menace known as the Joker emerges from his mysterious past, he
        'A thief who steals corporate secrets through the use of dream-sharing te
        'An ordinary hobbit named Frodo Baggins is recruited to carry the One Rin
        'Earth\'s mightiest heroes must come together and learn to fight as a tea
    ]
}

# Creating a DataFrame
movies_df = pd.DataFrame(data)

```

We will use TF-IDF (Term Frequency-Inverse Document Frequency) to convert the text data into a numerical format that we can use to calculate similarity.


```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity

# Combine the genre and description into a single feature
movies_df['combined_features'] = movies_df['genre'] + ' ' + movies_df['descriptio

# Create a TF-IDF Vectorizer
tfidf = TfidfVectorizer()
tfidf_matrix = tfidf.fit_transform(movies_df['combined_features'])

# Calculate the cosine similarity matrix
cosine_sim = cosine_similarity(tfidf_matrix, tfidf_matrix)

```

Now we will create a function that takes a movie title as input and returns a list of similar movies.

```

def recommend_movies(movie_title):
    # Get the index of the movie that matches the title
    idx = movies_df.index[movies_df['title'] == movie_title].tolist()[0]

    # Get the pairwise similarity scores of all movies with that movie
    sim_scores = list(enumerate(cosine_sim[idx]))

    # Sort the movies based on the similarity scores
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)

    # Get the scores of the 3 most similar movies
    sim_scores = sim_scores[1:4] # Exclude the first movie (itself)

    # Get the movie indices
    movie_indices = [i[0] for i in sim_scores]

    # Return the top 3 most similar movies
    return movies_df['title'].iloc[movie_indices]

# Example usage
recommended_movies = recommend_movies('The Matrix')
print(f"Movies recommended for 'The Matrix':\n{recommended_movies.tolist()}")

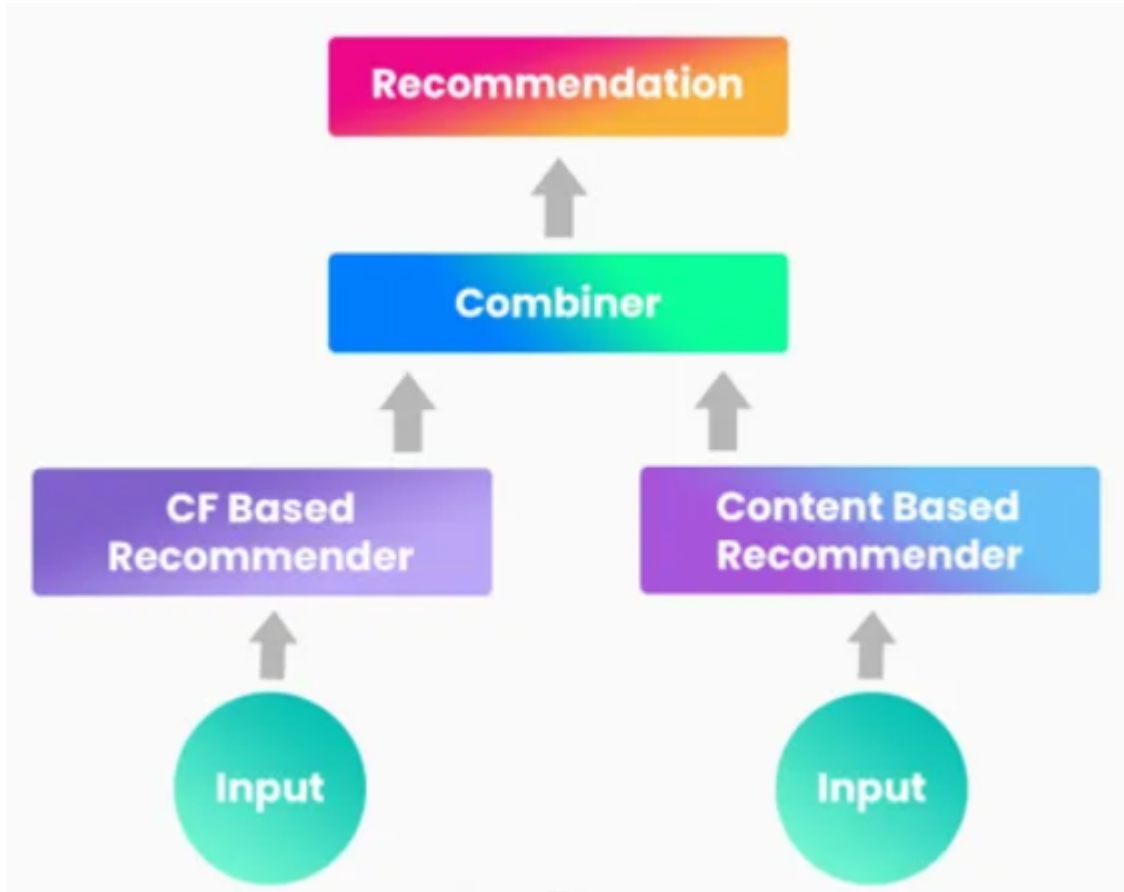
➡ Movies recommended for 'The Matrix':
['Inception', 'The Avengers', 'The Dark Knight']

```

✓ Hybrid Recommender Systems

- Hybrid recommender systems integrate various recommendation techniques to enhance the overall accuracy and robustness of recommendations.
- By combining the strengths of multiple approaches—such as collaborative filtering, content-based filtering, and demographic-based methods—hybrid systems can address the limitations of each individual technique.

- For instance, collaborative filtering may struggle with new users or items (the cold start problem), while content-based filtering can provide recommendations based solely on item attributes but may lack diversity.
- By merging these methods, hybrid systems can deliver more personalized and effective recommendations that adapt to diverse user preferences and behaviors.



✓ Key Components of Hybrid Recommender Systems

Hybrid recommendation systems combine both content-based and collaborative filtering techniques to provide more accurate and diverse recommendations.

1. Collaborative Filtering:

- Collaborative Filtering relies on the behavior and preferences of users to make recommendations. It can be divided into two main categories:
 - User-based Collaborative Filtering: This method recommends items based on the preferences of similar users. For example, if User A and User B have similar tastes, and User A liked a specific movie, that movie may be recommended to User B.
 - Item-based Collaborative Filtering: This approach recommends items that are similar to those the user has liked in the past. If a user enjoyed a particular

movie, the system will suggest other movies that similar users have also rated highly.

2. Content-Based Filtering:

- Content-Based Filtering recommends items by analyzing their attributes and the user's previous preferences. This technique creates a user profile based on the features of items they have liked or interacted with.
- The system then suggests new items that match the identified features. For instance, if a user enjoys action movies, the system will recommend other movies within the action genre or those that share similar characteristics, such as directors or actors.

✓ Types of Hybrid Recommender Systems

Hybrid recommender systems can be classified based on their integration strategies for combining different recommendation approaches. Here are the main categories of hybrid systems along with examples:

1. Weighted Hybrid:

- This method assigns different weights to various recommendation techniques and combines their outputs into a single recommendation list. The effectiveness of each approach is evaluated, and their scores are linearly combined.
- Example: A system may utilize both collaborative filtering and content-based filtering, assigning weights based on their predictive accuracy. For instance, if collaborative filtering is found to be 70% effective and content-based filtering is 30% effective, the final score for an item would reflect these weights.

2. Switching Hybrid:

- In a switching hybrid system, the choice of recommendation method is determined by specific conditions, such as user context or the availability of data. This allows the system to adapt to varying user scenarios.
- Example: For users with a rich interaction history, the system might primarily use collaborative filtering to generate recommendations. Conversely, for new users lacking history, the system would switch to content-based filtering to suggest items based on their attributes.

3. Mixed Hybrid:

- A mixed hybrid system combines the outputs of multiple recommendation techniques and presents them simultaneously to the user. This allows users to receive diverse recommendations from different sources.

- Example: The user interface might display items recommended through collaborative filtering alongside suggestions derived from content-based filtering, providing a richer set of options for the user.

4. Cascade Hybrid:

- In a cascade hybrid system, multiple recommendation methods are applied sequentially, where the output from one method serves as the input for another. This approach allows for more refined recommendations based on prior filtering.
- Example: A system may first filter items based on user demographics to create a relevant subset. Then, collaborative filtering could be applied to this filtered list to generate personalized recommendations from that smaller pool.

5. Feature Augmentation:

- This technique enhances one recommendation method by incorporating the output of another as additional features. This can improve the accuracy of the recommendations.
- Example: Collaborative filtering scores might be integrated as features in a content-based model, enriching the feature set and allowing for better predictions by leveraging insights from user behavior.

✓ Implementation Example

- We will use the following steps in our implementation:
 - Sample Data Preparation
 - Implementing Content-Based Filtering
 - Implementing Collaborative Filtering
 - Combining the Recommendations

Let's assume we have a dataset of movies with ratings and features (genres, descriptions).

```
import pandas as pd

# Sample data for movies with ratings
data = {
    'title': ['Movie A', 'Movie B', 'Movie C', 'Movie D', 'Movie E'],
    'genre': ['Action', 'Comedy', 'Action', 'Horror', 'Comedy'],
    'description': ['An action-packed adventure.',
                   'A funny comedy about friendship.',
                   'Another thrilling action movie.',
                   'A scary horror story.',
                   'A comedy that will make you laugh.'],
    'user_ratings': [5, 3, 4, 2, 5]
}

movies_df = pd.DataFrame(data)
```

Implementing Content-Based Filtering

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity

# Feature extraction using TF-IDF
tfidf = TfidfVectorizer()
tfidf_matrix = tfidf.fit_transform(movies_df['description'])
cosine_sim_content = cosine_similarity(tfidf_matrix, tfidf_matrix)
```

Implementing Content Based Filtering

```
# Function for content-based recommendations
def content_based_recommendation(movie_index):
    sim_scores = list(enumerate(cosine_sim_content[movie_index]))
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)
    return [movies_df['title'][i[0]] for i in sim_scores[1:3]] # Top 2 similar m
```

Implementing Collaborative Filtering

```
# Function for collaborative filtering based on ratings
def collaborative_filtering(movie_index):
    similar_movies = movies_df.loc[movies_df['user_ratings'] >= movies_df['user_r
    return similar_movies['title'].tolist()
```

Now, we'll combine the recommendations from both methods.

```
def hybrid_recommendation(movie_title):  
    movie_index = movies_df.index[movies_df['title'] == movie_title].tolist()[0]  
  
    # Get recommendations from both methods  
    content_recs = content_based_recommendation(movie_index)  
    collaborative_recs = collaborative_filtering(movie_index)  
  
    # Combine and remove duplicates  
    recommendations = list(set(content_recs + collaborative_recs))  
  
    return recommendations  
  
# Example usage  
recommended_movies = hybrid_recommendation('Movie A')  
print(f"Recommended movies for 'Movie A': {recommended_movies}")  
  
⇒ Recommended movies for 'Movie A': ['Movie C', 'Movie B', 'Movie E', 'Movie A']
```

✓ Handling Cold Start Problems

- The cold start problem in recommender systems arises when there is insufficient data to generate accurate recommendations, posing a significant challenge for new users, new items, or an entirely new system.
- To mitigate these challenges, various strategies can be employed, including user profiling to gather demographic information and preferences, content-based filtering to recommend items based on attributes, and hybrid approaches that combine multiple recommendation techniques.
- Additionally, utilizing popularity-based recommendations and active learning can help gather preference data quickly. By implementing these strategies, recommender systems can improve their effectiveness, even when faced with cold start scenarios.



✓ Types of Cold Start Problems

• New User Cold Start:

- The new user cold start problem arises when a user registers on a recommender system but has yet to provide any ratings or feedback on items. This lack of interaction makes it challenging for the system to infer their preferences and interests accurately.
- With no historical data available for the new user, the system struggles to deliver personalized recommendations.
- As a result, users may receive generic or irrelevant suggestions that do not align with their tastes. This can lead to frustration and disengagement from the platform, as users may feel that the system does not cater to their needs.
- To address this issue, systems often employ strategies like collecting demographic information during the onboarding process or utilizing popularity-based recommendations until sufficient data is gathered.

• New Item Cold Start:

- This problem occurs when a new item is introduced to the recommender system, but it has not yet garnered any ratings, reviews, or interactions from users. This absence of data makes it difficult to assess how well the new item fits within existing user preferences.
- Without ratings or user feedback, the system cannot evaluate the relevance or quality of the new item, leading to challenges in recommending it to users who might be interested.
- Consequently, new items may go unnoticed or underrepresented in recommendations, stunting their potential success. To mitigate this issue, content-

based filtering can be used, where recommendations are made based on the item's features and attributes, rather than relying solely on user interactions.

- Additionally, systems may encourage users to explore new items by promoting them as "trending" or "featured" options.

- **New System Cold Start:**

- A new system cold start problem arises when an entirely new recommender system is launched without any prior user-item interactions or historical data. This situation poses a unique challenge as the system begins without a foundation of user preferences or item characteristics to draw upon.
- The absence of historical data complicates the generation of initial recommendations, forcing the system to rely on generic methods that may not reflect the interests of its users.
- Users may encounter a lack of personalized content, leading to potential dissatisfaction and abandonment of the system. To address the new system cold start, strategies such as implementing user surveys to gather preference data, utilizing collaboration with partner platforms for initial data, or relying on content-based filtering techniques can be effective.
- Furthermore, leveraging machine learning algorithms that can quickly learn user preferences as they interact with the system can help to establish a more personalized experience over time.

✓ Challenges in Cold Start Scenarios

- **Data Sparsity:**

- In recommender systems, data sparsity occurs when there are very few interactions (e.g., ratings, clicks) between users and items, leading to sparse user-item matrices.
- This sparse data makes it difficult for algorithms, especially collaborative filtering, to compute meaningful similarities between users or items. With limited data points, the system may struggle to identify overlapping preferences, which can hinder accurate recommendations.

- **User Engagement:**

- When new users receive irrelevant or poorly matched recommendations due to data sparsity, their engagement with the system decreases.
- If users do not find the system useful early on, they may not interact further, creating a negative feedback loop where the system continues to lack user data.
- This disengagement exacerbates the cold start problem, preventing the system from learning more about user preferences over time.

- **Quality of Recommendations:**

- Sparse data negatively impacts the quality of recommendations. With insufficient data, the system may suggest items that are not personalized or relevant, frustrating users.
- Poor recommendation quality can lead to user dissatisfaction and, eventually, churn, as users abandon the platform for one that provides better suggestions.
- To improve recommendation quality in sparse data situations, systems may use hybrid methods (combining collaborative filtering with content-based filtering), popularity-based recommendations, or ask users for explicit preferences during onboarding.

✓ Strategies for Handling Cold Start Problems

1. Utilizing User Profiles

When a new user enters the system, there is little or no interaction data to work with. However, demographic information or self-reported preferences can provide a good starting point:

- **User Information:** During the registration process, ask users for basic demographic details, such as age, gender, and location, which can give insights into general preferences. For instance, younger users may have different preferences in movie genres compared to older users.
- **Onboarding Questionnaires:** To gather more specific data, systems can prompt new users to complete short surveys or questionnaires that ask about their tastes, preferences, or favorite items. This helps build an initial profile, allowing the system to make more relevant recommendations right away, without waiting for extensive interaction data to accumulate.

2. Content-Based Filtering

Content-based filtering focuses on the features of the items themselves, which makes it an excellent strategy for cold start situations, especially for new users or new items.

- **Feature Extraction:** In cases where new users or items are involved, item attributes like genre, keywords, or product descriptions are analyzed. For example, in a movie recommendation system, movies can be categorized by genre (e.g., action, comedy) or attributes such as actors, directors, or keywords. This data allows the system to recommend similar items that align with the user's preferences.
- **Content Similarity for New Items:** When a new item is added to the system and has no interaction history, its metadata (such as the product description, categories, or attributes) can be leveraged to match it with users who have shown a preference for similar items. This allows the system to recommend the new item to relevant users, even without user-item interaction data.

3. Hybrid Recommender Systems

Hybrid recommender systems combine multiple approaches—such as collaborative filtering and content-based filtering—to improve the overall recommendation quality, and they can significantly alleviate cold start issues.

- **Combining Techniques:** By integrating content-based filtering with collaborative filtering, the system can continue to make relevant recommendations even when one method alone might struggle due to insufficient data. For example, a new user with limited history might benefit from content-based recommendations, while users with rich interaction histories can benefit from collaborative filtering.
- **Fallback Mechanisms:** When collaborative filtering techniques fail due to lack of data (as with new users or items), hybrid systems can fall back on content-based filtering or demographic-based approaches to ensure recommendations are still provided. This adaptive approach allows the system to operate effectively in cold start scenarios.

4. Popularity-Based Recommendations

One simple but effective strategy for cold start users is to recommend popular or trending items, as these tend to have broader appeal.

- **Trending Items:** For new users who have little to no interaction history, recommending popular items or items that are currently trending can be a good starting point. These items typically have wide appeal and are more likely to engage new users.
- **Featured Items:** Highlighting new or featured items that are gaining attention can increase their visibility among new users, making it more likely that the items will receive early interactions and feedback, which can later be used for collaborative filtering.

5. Collaborative Filtering with Implicit Feedback

Instead of relying on explicit feedback such as ratings or reviews, systems can gather implicit feedback from user behavior, which is particularly useful in cold start situations.

- **Implicit Feedback Collection:** Systems can collect data based on user interactions, such as the number of clicks on an item, time spent viewing an item, or browsing patterns. These interactions can be used to infer user preferences even if the user hasn't provided explicit ratings.
- **Matrix Factorization with Implicit Feedback:** Advanced collaborative filtering techniques, such as matrix factorization (e.g., Alternating Least Squares or ALS), can be adapted to work with implicit data. This allows the system to build relationships between users and items even when only limited explicit interaction data is available. Implicit feedback helps the system capture user preferences more quickly, reducing the cold start impact.

6. Leveraging Social Networks

Social connections and shared preferences among users can provide valuable data to mitigate the cold start problem.

- **Social Data Integration:** By integrating data from social media or other platforms, systems can infer user preferences based on the behavior and preferences of their friends or social connections. For example, if a user is connected to friends who enjoy certain movies or products, the system can recommend similar items to that user.
- **Social Tagging and Sharing:** Implementing social tagging systems, where users tag items based on their personal interests, can also provide additional data points for the system to recommend items. These tags create an additional layer of metadata that can be used to align new items or users with the interests of others in the network.

✓ Example for Handling the Cold Start Problem

- We'll demonstrate a simple approach for handling the New User Cold Start problem by:
 - Recommending popular items when a user is new (no interaction history).
 - Gradually switching to collaborative filtering as the user interacts with the system.

We'll use a basic user-item interaction matrix with some missing entries to simulate the cold start scenario.

```
import pandas as pd
import numpy as np

# Sample user-item interaction data
data = {
    'user_id': [1, 2, 3, 4, 5],
    'movie_1': [5, 4, np.nan, 3, np.nan],
    'movie_2': [np.nan, 5, np.nan, np.nan, 4],
    'movie_3': [3, np.nan, 4, 2, 1],
    'movie_4': [np.nan, np.nan, 5, 4, np.nan],
    'movie_5': [np.nan, 4, np.nan, 5, 3]
}

# Create a DataFrame
ratings_df = pd.DataFrame(data)
ratings_df.set_index('user_id', inplace=True)
print(ratings_df)
```

	movie_1	movie_2	movie_3	movie_4	movie_5
user_id					
1	5.0	NaN	3.0	NaN	NaN
2	4.0	5.0	NaN	NaN	4.0
3	NaN	NaN	4.0	5.0	NaN
4	3.0	NaN	2.0	4.0	5.0
5	NaN	4.0	1.0	NaN	3.0

When a new user joins, we don't have any ratings from them. A simple approach is to recommend the most popular items (the ones with the highest average rating).

```
# Compute the mean rating for each movie (ignoring NaN values)
mean_ratings = ratings_df.mean()

# Recommend the top 3 popular movies based on average rating
top_popular_movies = mean_ratings.sort_values(ascending=False).head(3)

print("Top 3 popular movies for a new user:")
print(top_popular_movies)
```

```
⇒ Top 3 popular movies for a new user:
movie_2    4.5
movie_4    4.5
movie_1    4.0
dtype: float64
```

- Here, we see that Movie 4, Movie 1, and Movie 5 are recommended to new users based on the average ratings across all users.
- This approach works well when we have no prior information about a new user's preferences, and it leverages the popularity of movies to provide reasonable recommendations.

Once the new user interacts with a few items, we can switch to Collaborative Filtering. We'll use User-User Collaborative Filtering in this case. For simplicity, we'll use Cosine Similarity to find similar users and recommend movies based on their ratings.

Compute User Similarity

```
from sklearn.metrics.pairwise import cosine_similarity

# Fill NaN values with 0 for similarity calculation
user_ratings_filled = ratings_df.fillna(0)

# Compute the cosine similarity between users
user_similarity = cosine_similarity(user_ratings_filled)
user_similarity_df = pd.DataFrame(user_similarity, index=ratings_df.index, column

print("User Similarity Matrix:")
print(user_similarity_df)
```

```
⇒ User Similarity Matrix:
user_id      1      2      3      4      5
user_id
1      1.000000  0.454311  0.321403  0.490098  0.100901
2      0.454311  1.000000  0.000000  0.576787  0.831239
3      0.321403  0.000000  1.000000  0.595072  0.122513
4      0.490098  0.576787  0.595072  1.000000  0.453696
5      0.100901  0.831239  0.122513  0.453696  1.000000
```

- This matrix is key to making recommendations using collaborative filtering.

- The more similar two users are, the more likely they will enjoy the same items.

We'll assume that the user has rated at least one movie (thus transitioning from the cold start phase). We'll use the ratings of similar users to recommend new movies.

```
def recommend_for_existing_user(user_id, n_recommendations=3):  
    # Get similarity scores for the input user  
    similar_users = user_similarity_df[user_id].sort_values(ascending=False)
```